Linear Regression Model for House Price Prediction

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score

import matplotlib.pyplot as plt

import seaborn as sns


# Load the dataset

train_data = pd.read_csv('/mnt/data/train.csv')


# Select relevant features

features = ['GrLivArea', 'BedroomAbvGr', 'FullBath']

X = train_data[features]

y = train_data['SalePrice']


# Split the data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Create and train the model

model = LinearRegression()

model.fit(X_train, y_train)


# Make predictions
```

```
y_pred = model.predict(X_test)


# Evaluate the model

mse = mean_squared_error(y_test, y_pred)

r2 = r2_score(y_test, y_pred)


# Visualize the results

plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred, color='blue')

plt.plot([y.min(), y.max()], [y.min(), y.max()], 'k--', lw=3)

plt.xlabel('Actual')

plt.ylabel('Predicted')

plt.title('Actual vs Predicted Sale Price')

plt.show()
```

Model Interpretation

Linear Regression Model Evaluation:

- Mean Squared Error (MSE): 2806426667.247853

- R-squared (R^2) Score: 0.6341189942328371

The R^2 score indicates that 63.41% of the variance in the sale price can be explained by the model using the given features. The scatter plot visualizes the relationship between the actual and predicted sale prices, showing how well the model's predictions align with the actual values.

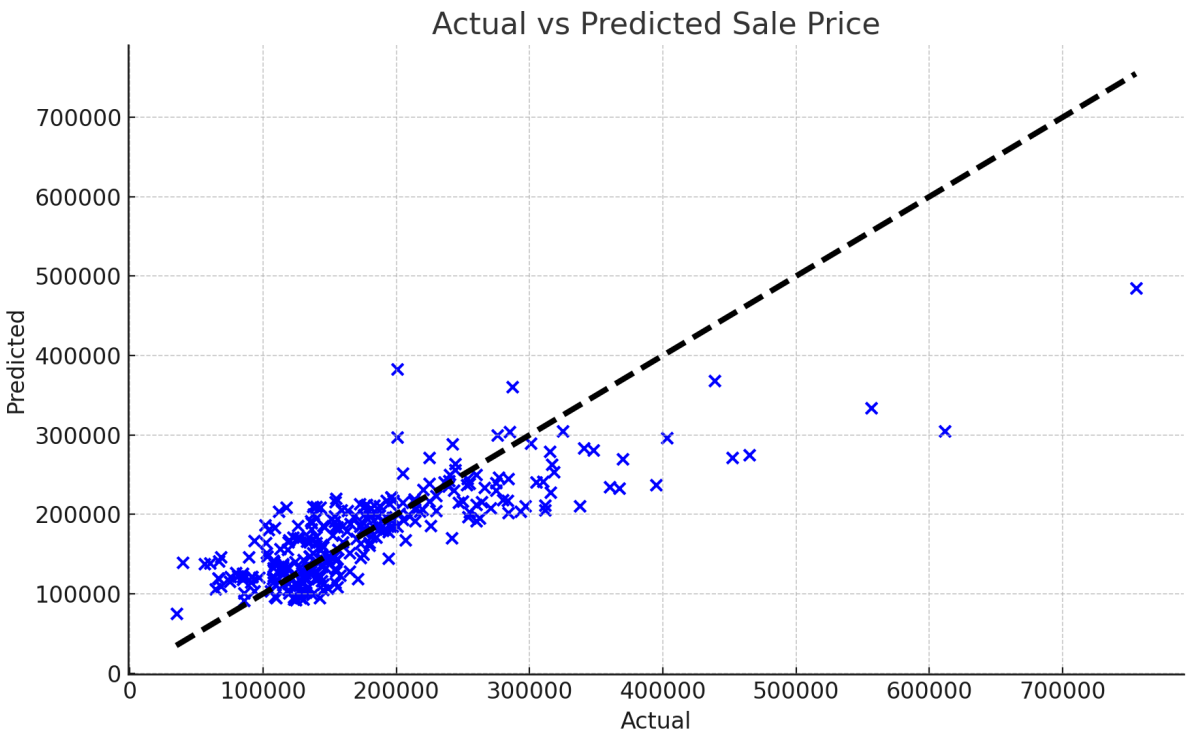
Actual vs Predicted Sale Price
Predicted
Actual
0
100000
200000
300000
400000
500000
600000
700000